

Quantum Integer Programming

47-779

Novel Ising Solvers

Carnegie Mellon University

Tepper School of Business

William Larimer Mellon, Founder



Agenda

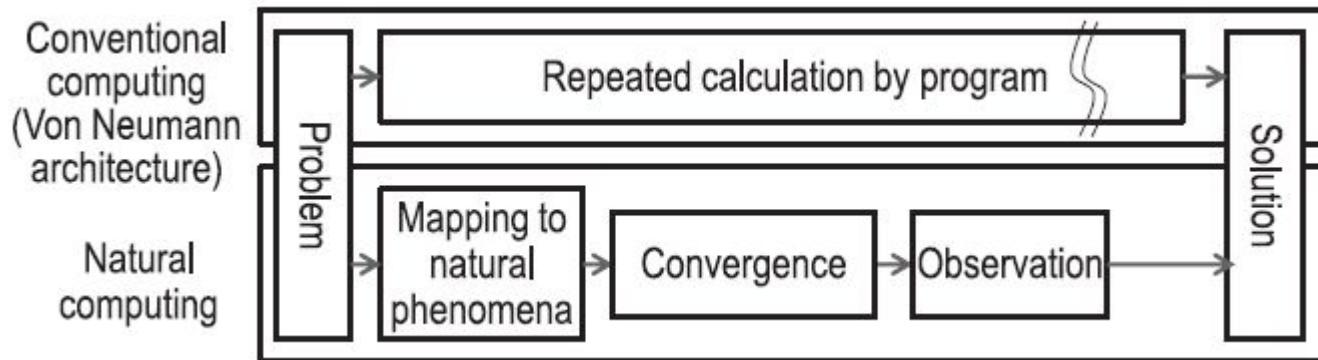
- Conventional vs. Natural Computing
- Solving the 2D regular Ising Problem
 - Graphic Processing Units
 - Tensor Processing Units
 - Field-programmable gate arrays
- Solving general Ising models
 - Graphic Processing Units
 - Simulated Bifurcation Machine
 - CMOS
 - Digital Annealers

Carnegie Mellon University

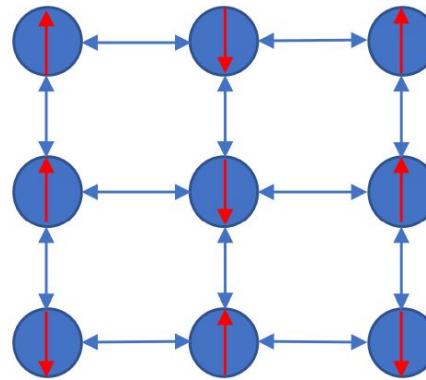
Tepper School of Business

William Larimer Mellon, Founder

Conventional (Von Neumann) vs. Natural Computing



2D Ising model - Simple yet interesting

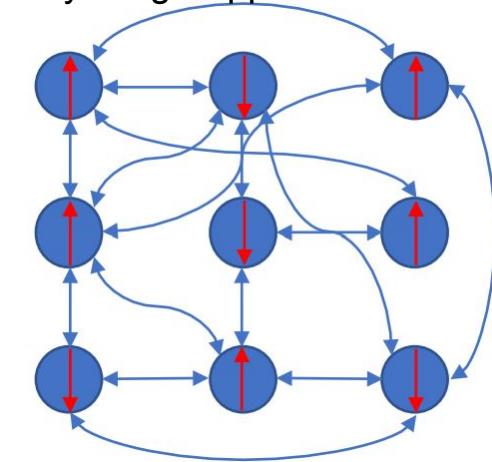


Ising
Models

Main concern: How to parallelize
Monte Carlo Simulations

[1] A 20k-Spin Ising Chip to Solve Combinatorial Optimization Problems With CMOS Annealing. Yamaoka, Yoshimura, Hayashi, Okuyama, Aoki, and Mizuno

Arbitrary Ising - Applicable but hard!



Main concern: How to actually solve
NP-Hard Problem

[2] <https://arxiv.org/pdf/1807.10750.pdf>

Carnegie Mellon University

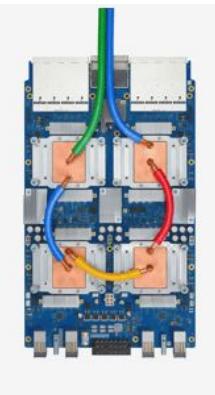
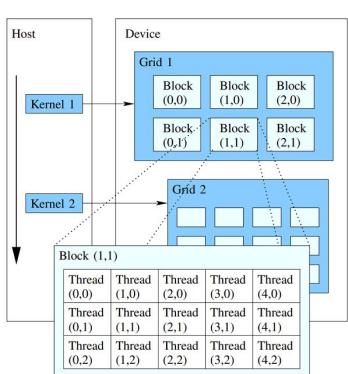
Tepper School of Business

William Larimer Mellon, Founder

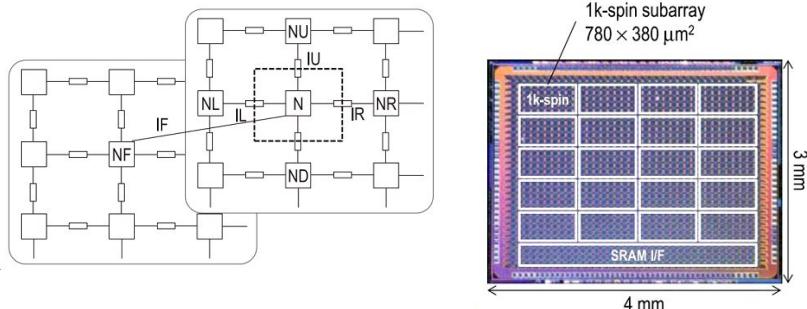


Specialized Hardware for Ising/QUBO

GPUs and TPUs



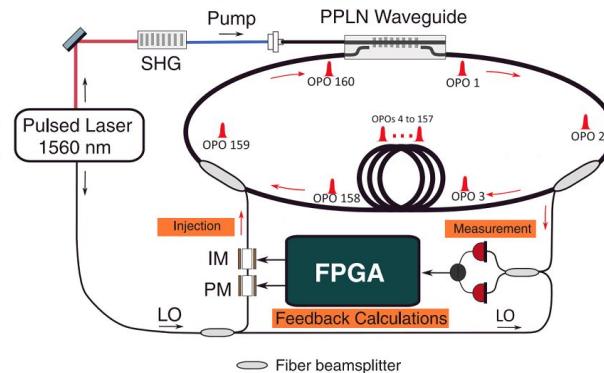
Complementary metal-oxide semiconductors (CMOS)



Digital annealers



Oscillator Based Computing



[1]<https://arxiv.org/pdf/1807.10750.pdf>

[2]<https://arxiv.org/pdf/1903.11714.pdf>

[3]<https://arxiv.org/pdf/1806.08815.pdf>

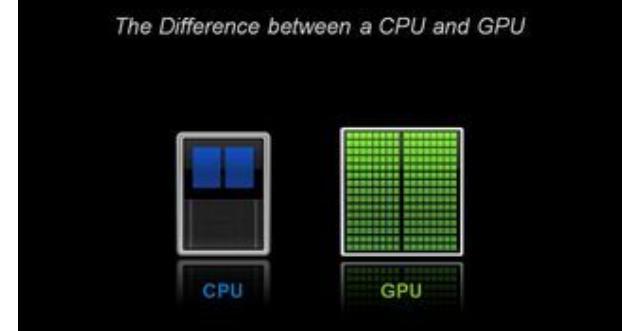
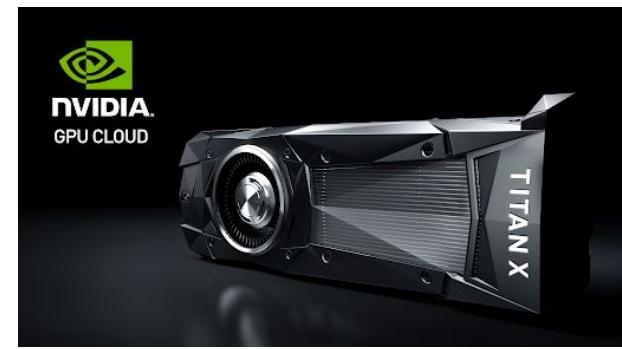
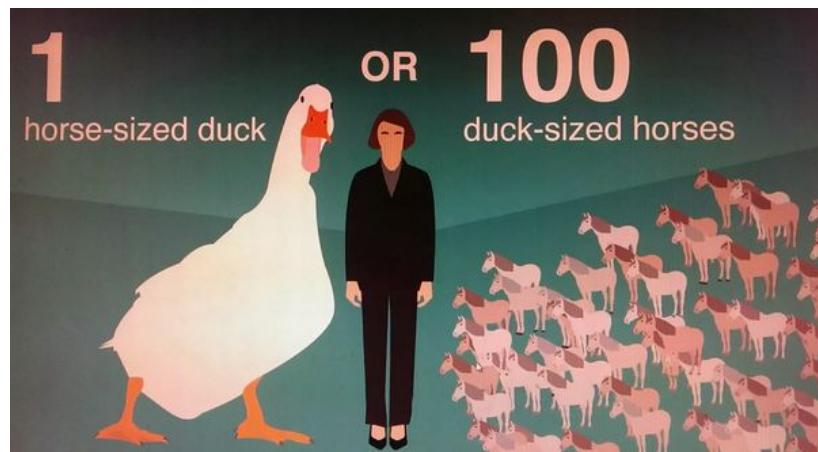
[4]<https://spectrum.ieee.org/tech-talk/mos-digital-annealer-produces-quantum-computer-speeds>

[5]<https://science.sciencemag.org/content/sci/354/6312/614.full.pdf>

Graphical Processing Units (GPU)

CPU vs GPU

CPU	GPU
Central Processing Unit	Graphics Processing Unit
Several cores	Many cores
Low latency	High throughput
Good for serial processing	Good for parallel processing
Can do a handful of operations at once	Can do thousands of operations at once



Specialized, electronic circuit designed to rapidly manipulate and alter memory to accelerate the creation of images.... Their highly parallel structure makes them more efficient than general-purpose central processing units (CPUs) for algorithms that process large blocks of data in parallel.

1. <https://blogs.nvidia.com/blog/2009/12/16/whats-the-difference-between-a-cpu-and-a-gpu/>
2. <https://www.quora.com/Would-you-rather-fight-100-duck-sized-horses-or-one-horse-sized-duck>
3. https://en.wikipedia.org/wiki/Graphics_processing_unit

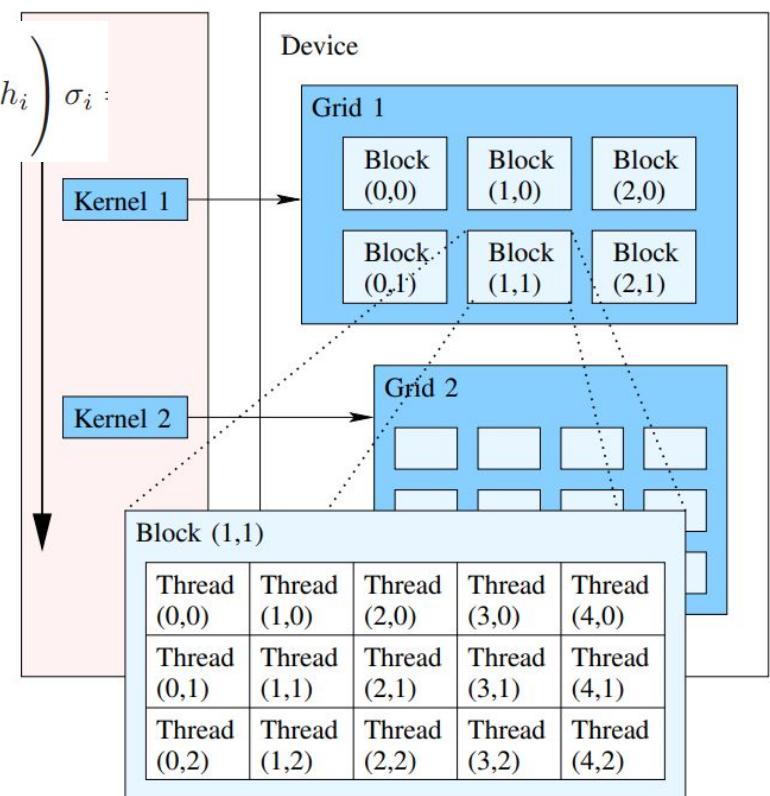
Graphical Processing Units

Algorithm 1 Modified Ising annealing algorithm

```
1: input: ( $M, N, S$ )
2: initialize all  $\sigma_i$  in  $S$ 
3: for sweep-id in  $\{1, 2, \dots, M\}$  do
4:   for  $\sigma_i$  in  $S$  do
5:      $\sigma_i \leftarrow \text{argmin}(H(\sigma_i))$  based on  $H_i(\sigma_i) = \left( - \sum_j J_{i,j} \sigma_j - h_i \right)$ 
6:   end for
7:   randomly choose and flip  $N$  spin glasses in  $S$ 
8:   decrease  $N$ 
9: end for
```

Algorithm 2 GPU Simulated Annealing method for Ising model

```
input: ( $F_p, S$ )
initialize ALL  $\sigma_i$  in  $S$ 
while  $F_p > 0$  do
  for all  $\sigma_i \in S$  in parallel do
     $\sigma_i \leftarrow \text{argmin}(H(\sigma_i))$ 
    flip  $\sigma_i$  with probability  $F_p$ 
  end for
  reduce  $F_p$ 
end while
```



1. <https://arxiv.org/pdf/1807.10750.pdf>

William Larimer Mellon, Founder

GPU for 2D Ising Models

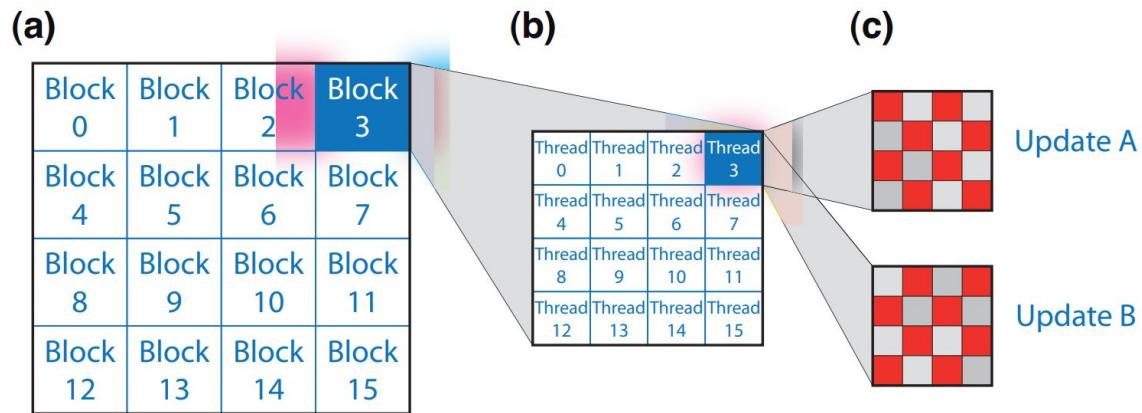


Figure 1: (Color online) The spin lattice is processed by a variable number of blocks (a), where each block runs a variable number of threads (b). The threads update the spin lattice in two steps, A and B, using two kernel invocations (c).

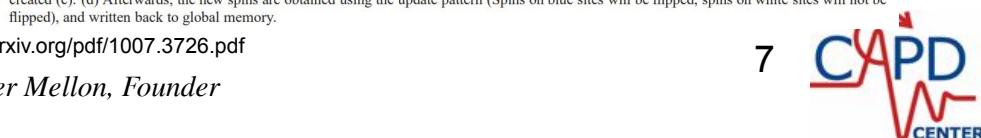
Perform the Ising Update via easily parallelizable operations.

	Spinflips per μs	Relative speed
CPU simple	26.6	0.11
CPU multi-spin coding	226.7	1.00
shared memory	4415.8	19.50
shared memory (Fermi)	8038.2	35.46
multi-spin unmodified	3307.2	14.60
multi-spin coding on the fly	5175.8	22.80
multi-spin coding linear	7977.4	35.20

Carnegie Mellon University
Tepper School of Business

1. <https://arxiv.org/pdf/1007.3726.pdf>
William Larimer Mellon, Founder

Figure 2: (Color online) (a) The way a kernel processes a 4×4 meta-spin. (b) Spins are extracted into shared memory and an update pattern is created (c). (d) Afterwards, the new spins are obtained using the update pattern (Spins on blue sites will be flipped, spins on white sites will not be flipped), and written back to global memory.



Correctness of GPU's results

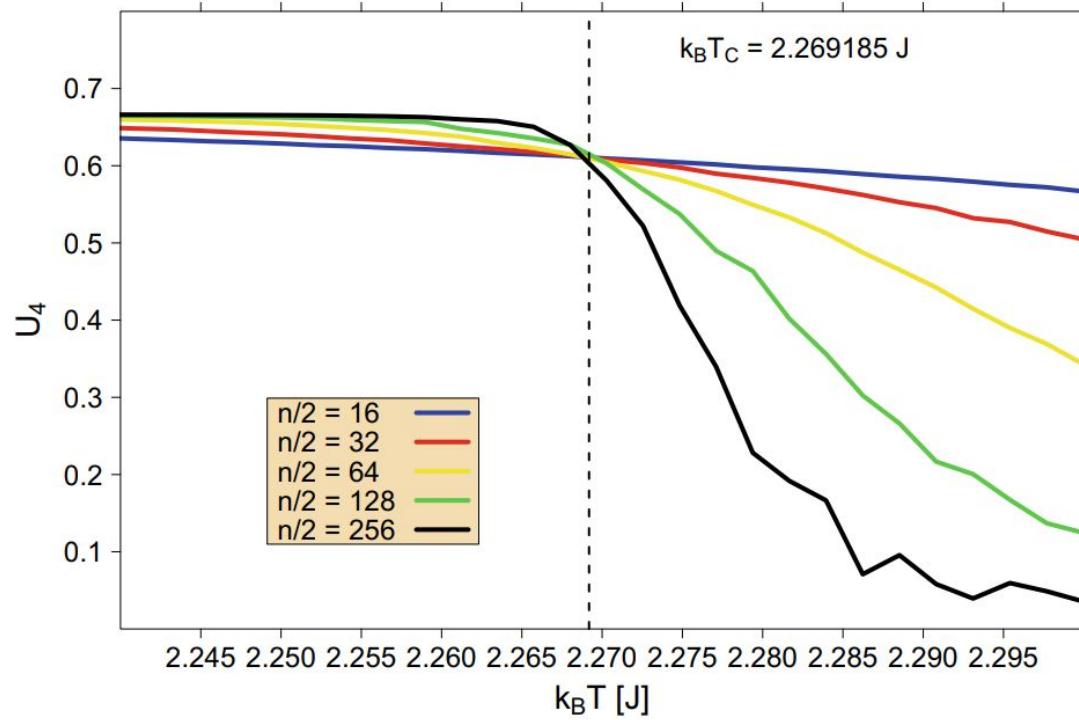


Fig. 5. Binder cumulant U_4 in dependence of $k_B T$ for various numbers n of spins per row and column of the two dimensional square lattice Ising model. $n/2$ corresponds to the involved number of threads per block on the GPU implementation. The curves of the Binder cumulants for various system sizes $N = n^2$ cross almost perfectly at the critical temperature derived by Onsager [3], which is shown additionally as a dashed line. In each temperature step, the average was taken over 10^7 measurements.

It (sort of) matches Onsager analytical prediction!

Parallelizing GPUs

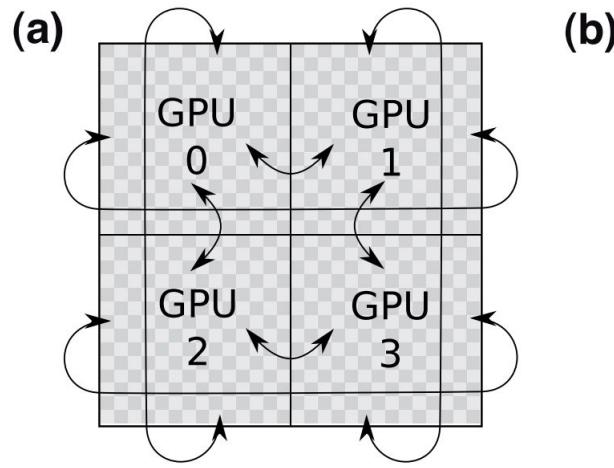
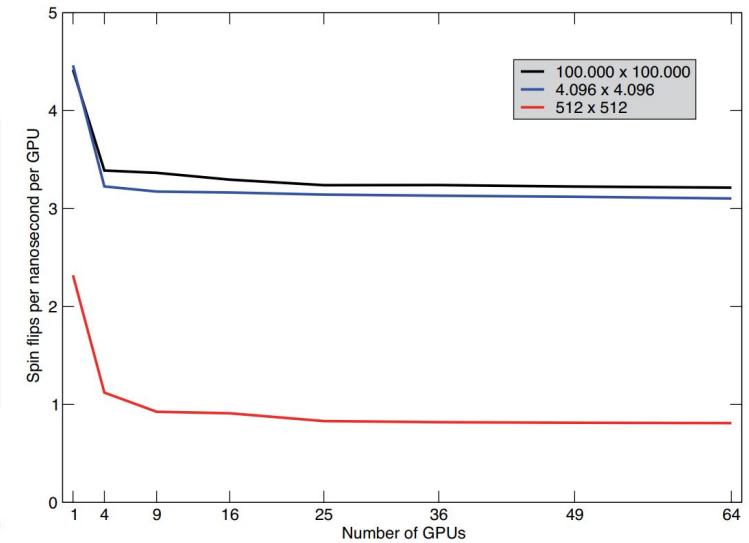
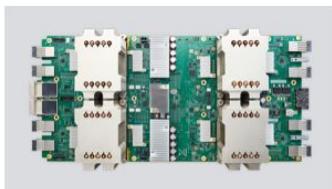


Figure 4: (Color online) (a) Each GPU processes a “meta-spin” lattice of size $N = n^2$. The lattices are aligned on a super-lattice, and the outer borders are connected via periodic boundary conditions. In this example, 4 GPUs process a system of $2^2 \cdot N$ spins. (b) A meta-spin update needs the 4 nearest neighbor meta-spins. On the borders of a lattice, each GPU needs the spin information of the neighboring lattices. The border information has to be passed between the GPUs. In our implementation this is done by using 8 neighbor arrays.

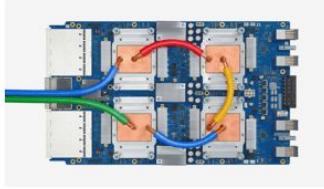


Using 64 GPUs performance of 206 spinflips per nanosecond can be achieved on a $800,000 \times 800,000$ lattice

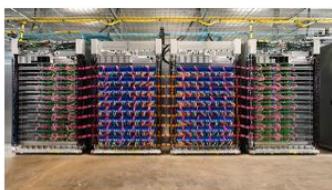
Tensor Processing Units (TPU)



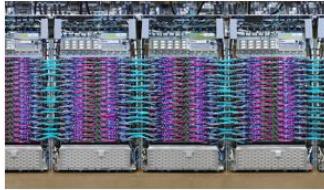
Cloud TPU v2
180 teraflops
64 GB High Bandwidth Memory (HBM)



Cloud TPU v3
420 teraflops
128 GB HBM



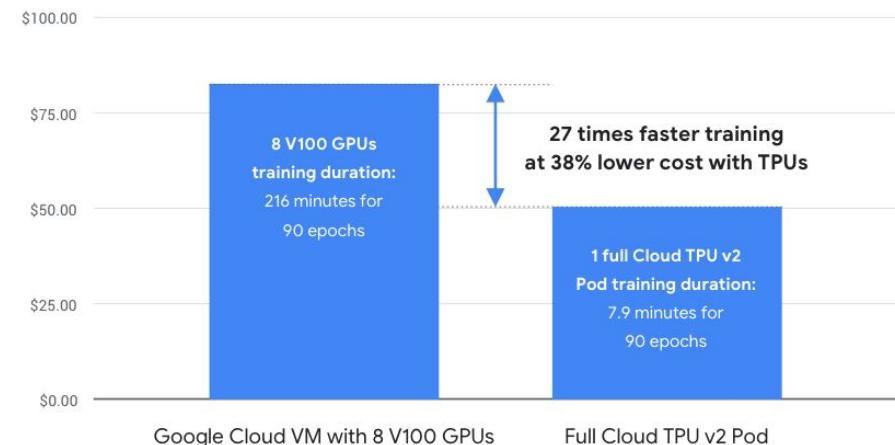
Cloud TPU v2 Pod
11.5 petaflops
4 TB HBM
2-D toroidal mesh network



Cloud TPU v3 Pod
100+ petaflops
32 TB HBM
2-D toroidal mesh network

Machine learning performance and benchmarks

ResNet-50 Training Cost Comparison



Tensor Processing Unit (TPU) is an AI accelerator application-specific integrated circuit (ASIC) developed by [Google](#) specifically for neural network machine learning, particularly using Google's own [TensorFlow](#) software.

1. https://en.wikipedia.org/wiki/Tensor_Processing_Unit#/media/File:Tensor_Processing_Unit_3.0.jpg
2. <https://cloud.google.com/tpu>

William Larimer Mellon, Founder

TPU for 2D Ising

Checkerboard Algorithm: break lattice in sublattices and group equal spins to easily operate on them

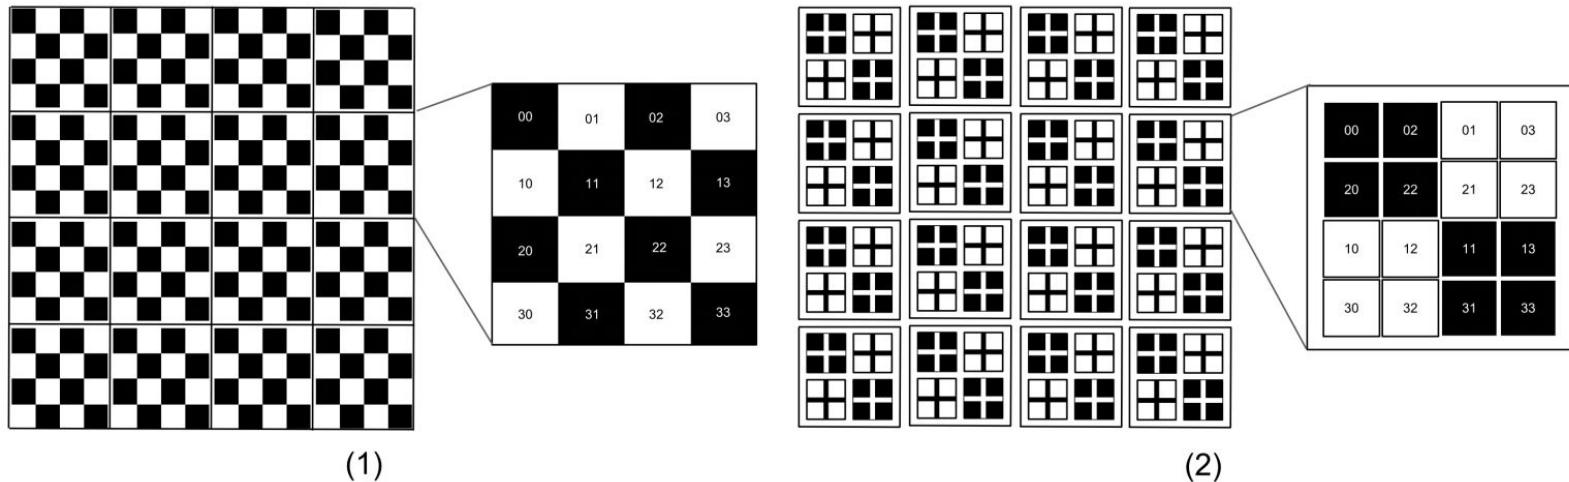
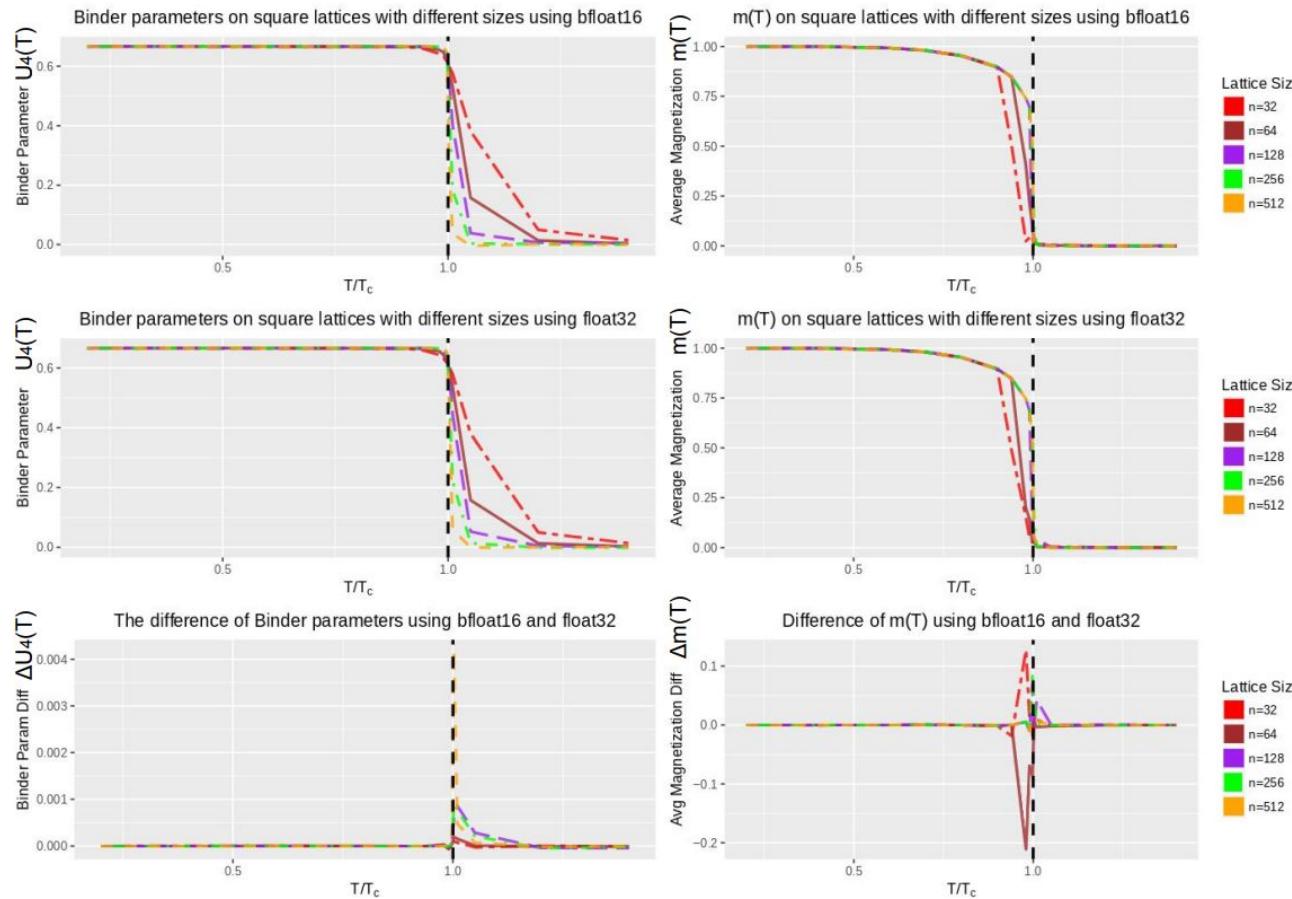


Figure 3: A 2-d checkerboard: (1) Original checkerboard: on the left, the 16×16 board is split into a 4×4 grid of 4×4 sub-lattices, i.e., it is represented by a $[4, 4, 4, 4]$ tensor, where $[l, k, :, :]$ is the sub-lattice at $[l, k]$ of the grid; on the right, the sub-lattice is zoomed in and the indices of its spin sites are shown; (2) Reorganized checkerboard: one the left, each 4×4 sub-lattice is reorganized by 4 “compact” 2×2 sub-lattices; on the right, 4 “compact” 2×2 sub-lattices are zoomed in and their original indices from the 4×4 sub-lattice are shown. In general, such alternate coloring of black and white can be extended to lattices with any dimensions.

1. <https://arxiv.org/pdf/1903.11714.pdf>

Correctness of TPU's results

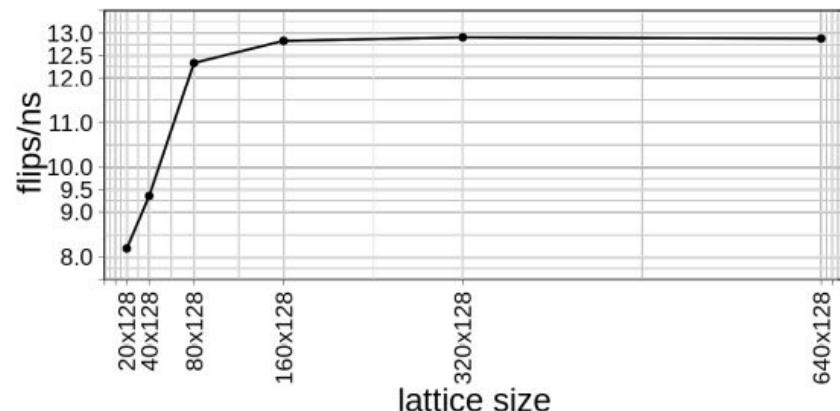


It (sort of) matches
Onsager analytical
prediction!

1. <https://arxiv.org/pdf/1903.11714.pdf>

Efficiency of TPU cluster

lattice size n^2	(flips/ns)	(nJ/flip)
$(20 \times 128)^2$	8.1920	12.2070
$(40 \times 128)^2$	9.3623	10.6811
$(80 \times 128)^2$	12.3362	8.1062
$(160 \times 128)^2$	12.8266	7.7963
$(320 \times 128)^2$	12.9056	7.7486
$(640 \times 128)^2$	12.8783	7.7650
GPU in [23, 3]	7.9774	—
Nvidia Tesla V100	11.3704	21.9869
FPGA in [20]	614.4	—



Better performance and less energy consumption than Nvidia GPUs, until... (next slide)

Really far from Field-programmable gate array (FPGA)! (a couple slides more)

Code available in Github and replicable results (with a Google Cloud account)

https://github.com/google-research/google-research/blob/master/simulation_research/ising_model/ising_mcmc_tpu.ipynb

1. <https://arxiv.org/pdf/1903.11714.pdf>

Nvidia's Rebuttal!

lattice size	flip/ns
$(1 \times 2048)^2$	231.09
$(2 \times 2048)^2$	318.95
$(4 \times 2048)^2$	379.27
$(8 \times 2048)^2$	411.65
$(16 \times 2048)^2$	420.44
$(32 \times 2048)^2$	420.77
$(64 \times 2048)^2$	418.23
$(123 \times 2048)^2$	417.53
1 TPUv3 core in [7]	12.91
32 TPUv3 cores in [7]	336.01
FPGA (1024^2) [8]	614.40 [†]

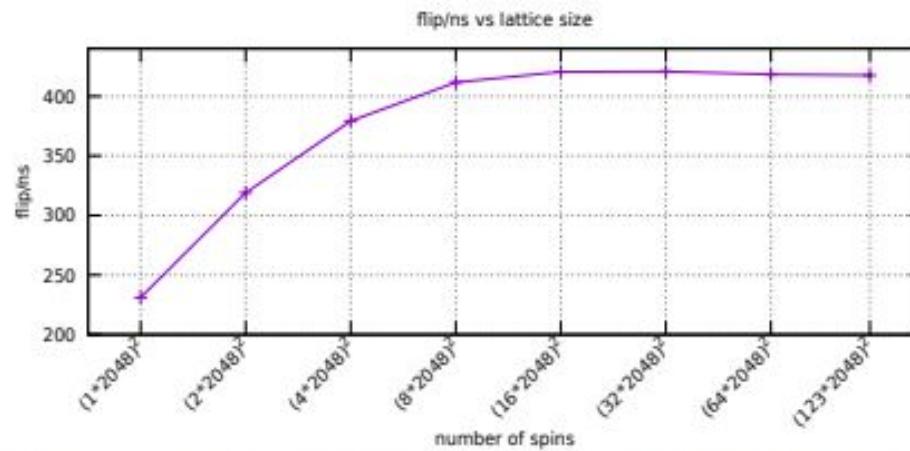


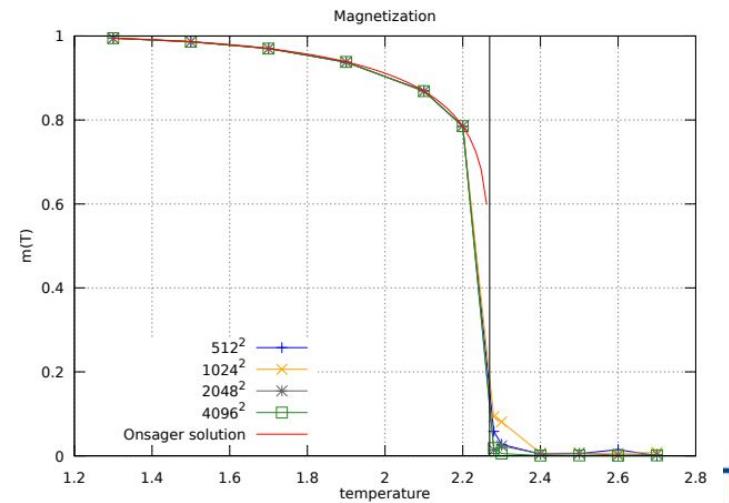
Table 2: Flips per nanosecond obtained by the optimized multi-spin code on a single Tesla V100-SXM card with different lattice sizes, requiring an amount of memory ranging from 2MB to 30GB. For comparison purposes, the table also reports the best timings with 1 and 32 TPUv3 cores from [7], and with 1 FPGA from [8].

Better performance than TPUs

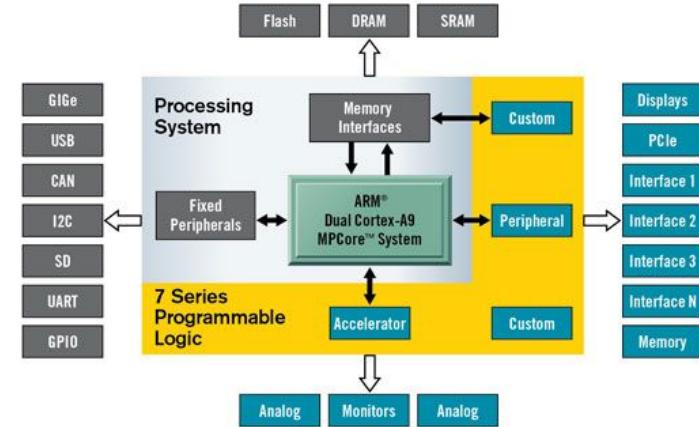
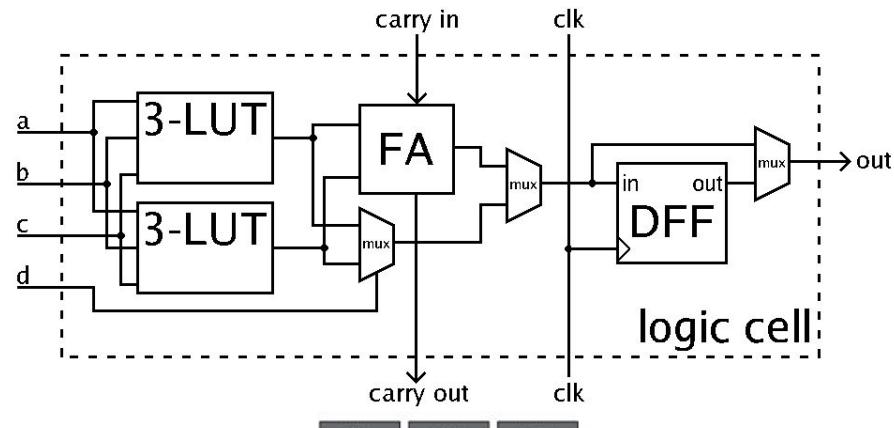
Still far from Field-programmable gate array (FPGA)! (next slide)

Code available in Github and replicable results <https://github.com/NVIDIA/ising-gpu>

And of course they compare against Onsager



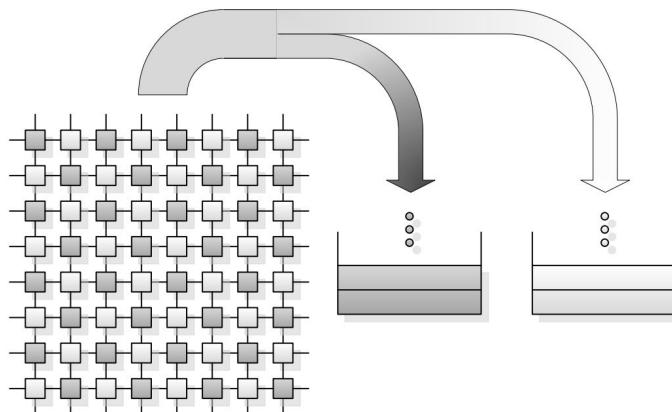
Field-programmable gate array (FPGA)



A **field-programmable gate array (FPGA)** is an **integrated circuit** designed to be configured by a customer or a designer after manufacturing – hence the term "**field-programmable**"... Circuit diagrams were previously used to specify the configuration, but this is increasingly rare due to the advent of **electronic design automation** tools.

1. https://en.wikipedia.org/wiki/Field-programmable_gate_array
2. <https://arxiv.org/pdf/1602.03016.pdf>

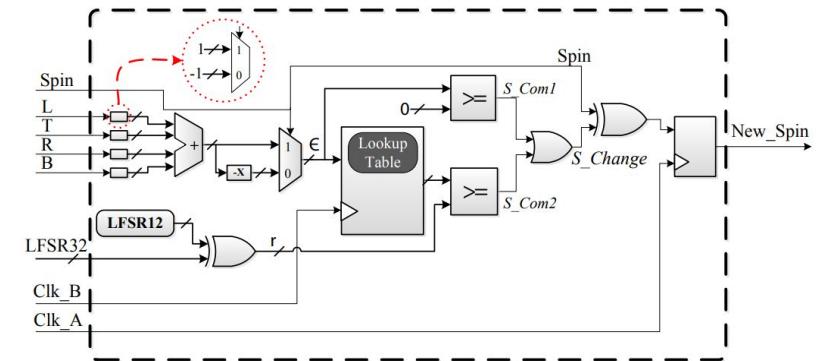
FPGA for 2D Ising Models



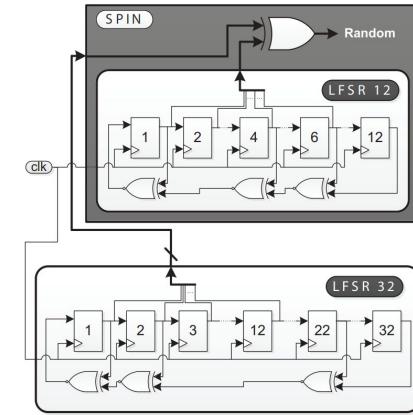
Checkerboard Algorithm diagram

Platform	# updated spins	Ratio
CPU	62	1
Single GPU	7977	129
Previous FPGA	94127	1518
64 GPUs	206000	3322
Our FPGA	614400	9909

Number of spinflips per microsecond for the 1024x1024 lattice



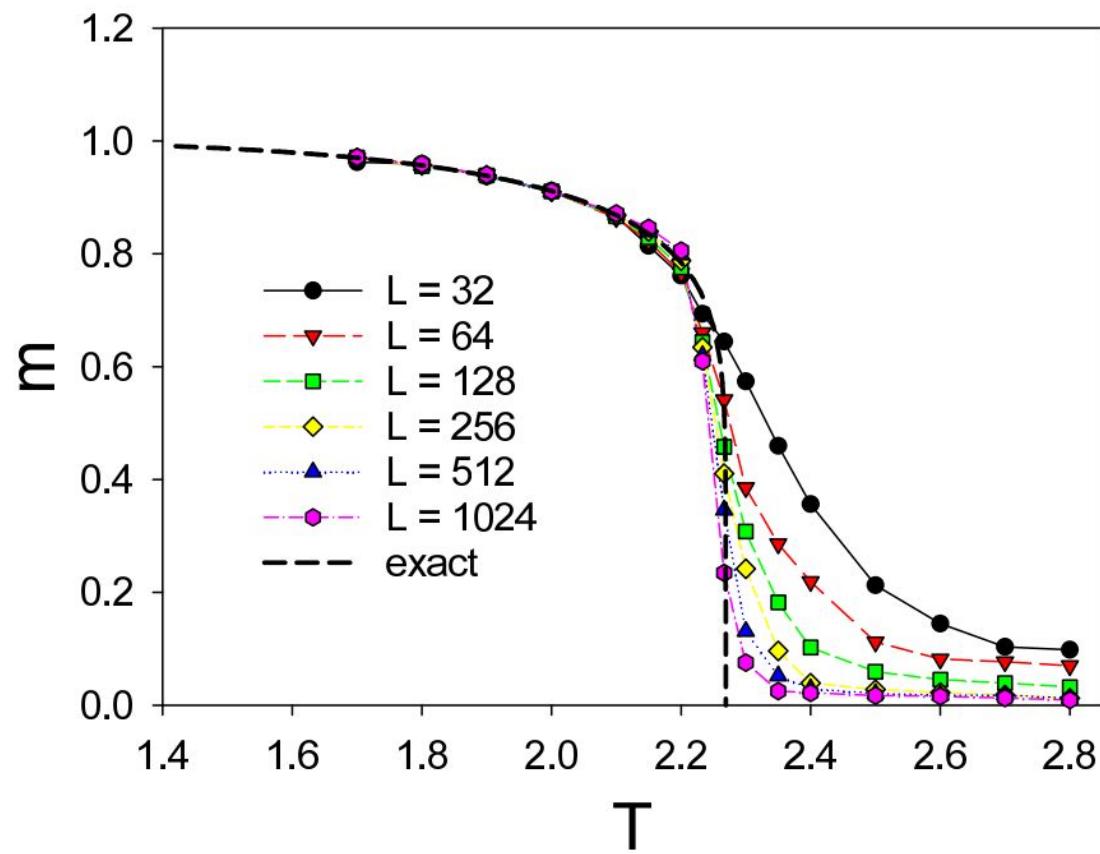
Circuit Diagram Single Spin



Circuit Diagram Random Number Generation

1. <https://arxiv.org/pdf/1602.03016.pdf>

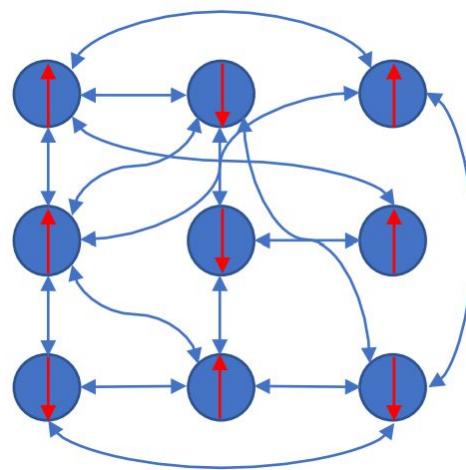
Correctness of FPGA's results



It (sort of) matches Onsager analytical prediction!

1. <https://arxiv.org/pdf/1602.03016.pdf>

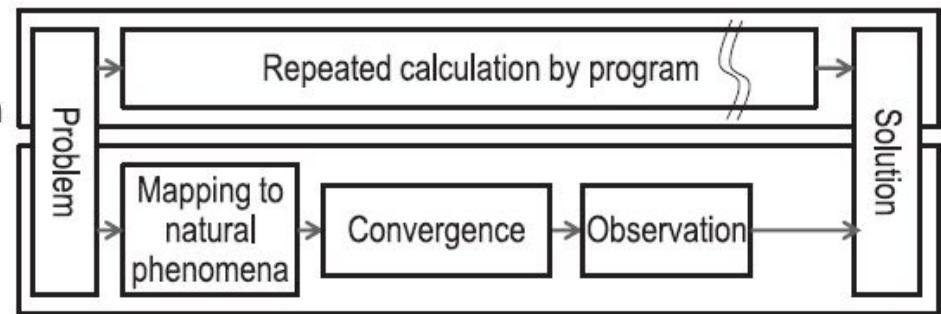
Working with general Ising Models



Main concern: How to actually solve NP-Hard Problem?

Conventional computing
(Von Neumann architecture)

Natural computing



Using a natural computing approach you would ideally use Adiabatic Quantum Computing, and realistically Quantum Annealing

Schrödinger equation

$$i\eta \frac{d}{dt}|\psi\rangle = H|\psi\rangle \quad H(\sigma_1, \sigma_2, \dots, \sigma_n) = -\frac{1}{2} \sum_i \sum_j J_{i,j} \sigma_i \sigma_j + \sum_i h_i \sigma_i$$

One cannot efficiently solve this equation using classical computers
(if so, why would we need quantum computers after all!)

The issue then relies on (classically) Simulating Quantum Annealing

[1] A 20k-Spin Ising Chip to Solve Combinatorial Optimization Problems With CMOS Annealing. Yamaoka, Yoshimura, Hayashi, Okuyama, Aoki, and Mizuno

Carnegie Mellon University

Tepper School of Business

William Larimer Mellon, Founder

[2] <https://arxiv.org/pdf/1807.10750.pdf>

Graphical Processing Units

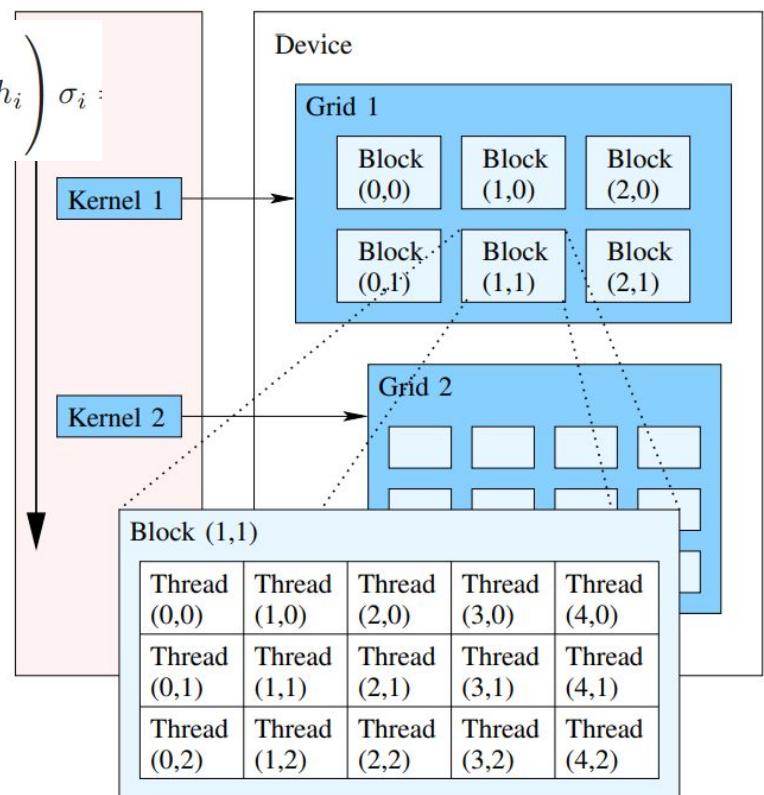
Algorithm 1 Modified Ising annealing algorithm

```
1: input:  $(M, N, \mathbf{S})$ 
2: initialize all  $\sigma_i$  in  $\mathbf{S}$ 
3: for sweep-id in  $\{1, 2, \dots, M\}$  do
4:   for  $\sigma_i$  in  $\mathbf{S}$  do
5:      $\sigma_i \leftarrow \text{argmin}(H(\sigma_i))$  based on  $H_i(\sigma_i) = \left( - \sum_j J_{i,j} \sigma_j - h_i \right)$ 
6:   end for
7:   randomly choose and flip  $N$  spin glasses in  $\mathbf{S}$ 
8:   decrease  $N$ 
9: end for
```

Algorithm 2 GPU Simulated Annealing method for Ising model

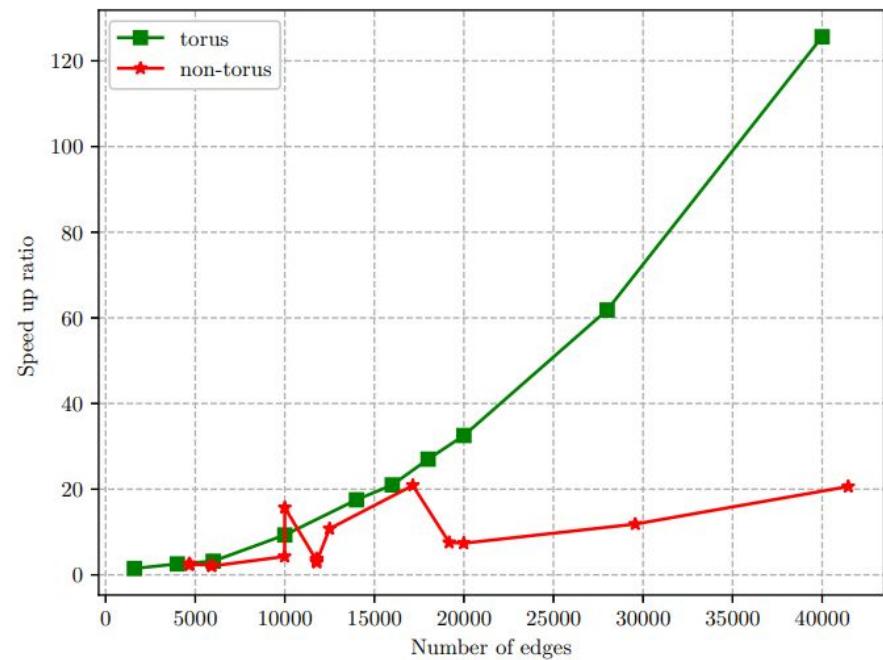
```
input:  $(F_p, \mathbf{S})$ 
initialize ALL  $\sigma_i$  in  $\mathbf{S}$ 
while  $F_p > 0$  do
  for all  $\sigma_i \in \mathbf{S}$  in parallel do
     $\sigma_i \leftarrow \text{argmin}(H(\sigma_i))$ 
    flip  $\sigma_i$  with probability  $F_p$ 
  end for
  reduce  $F_p$ 
end while
```

“One may notice that since each spin glass may have a different number of neighbors, then the threads will not be perfectly load balanced.”

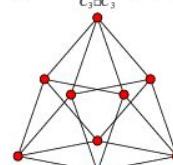


GPU Performance

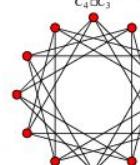
# edges	CPLEX cut	Days	Seconds
		GPU cut	(%accuracy)
9999	9473		8884 (93.78%)
14999	13357		12776(95.65%)
24998	20206		19981(98.88%)
49995	35248		36228(100.29%)
39998	33605		32914(97.94%)
59997	46371		46510(100.29%)
99995	70566		72009(102.04%)
199990	128448		131930(102.71%)
249995	176556		179391(101.60%)
374993	248505		255078(102.64%)
626988	392912		400540(101.94%)
1249975	741709		751050(101.25%)



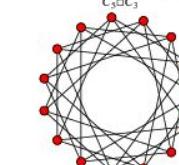
generalized quadrangle (2,1)



12-circulant graph (3,4)



15-circulant graph (3,5)



<https://arxiv.org/pdf/1807.10750.pdf>
<https://mathworld.wolfram.com/TorusGridGraph.html>

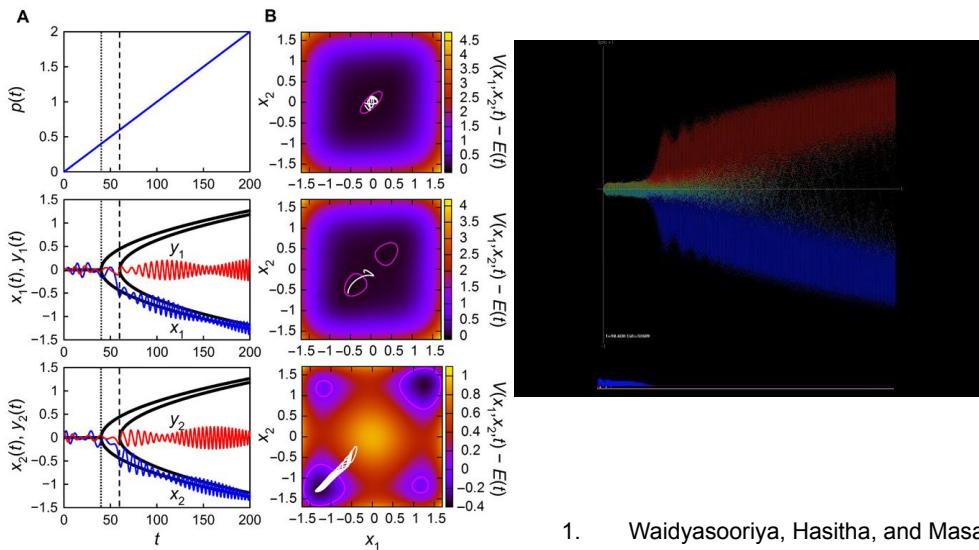
Simulated Bifurcation Machine

“The method in [previous slide] ignores the data dependencies to implement parallel computation on fully connected spin models. Since the modified algorithm in [previous slide] does not follow the mathematical model that the Quantum Monte Carlo is based on, the output of the simulation could deviate from the optimum.”

How can we efficiently simulate quantum annealing?

We can take a classical approximation

Equations model the bifurcation (Anil's lecture)



$$H_{SB}(\vec{x}, \vec{y}, t) = \sum_{i=1}^N \frac{\Delta}{2} y_i^2 + V(\vec{x}, t)$$

$$V(\vec{x}, t) = \sum_{i=1}^N \left[\frac{K}{4} x_i^4 + \frac{\Delta - p(t)}{2} x_i^2 \right] - \frac{\xi_0}{2} \sum_{i=1}^N \sum_{j=1}^N J_{i,j} x_i x_j$$

$$\frac{\partial x_i}{\partial t} = \frac{\partial H_{SB}}{\partial y_i} = \Delta y_i$$

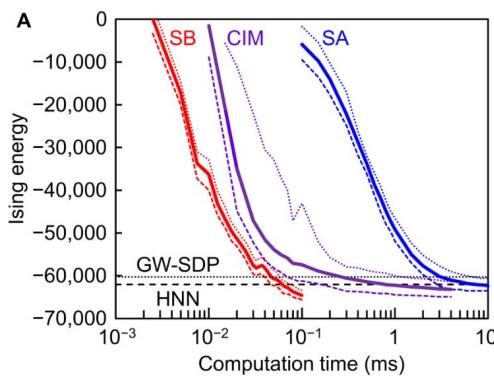
$$\frac{\partial y_i}{\partial t} = -\frac{\partial H_{SB}}{\partial x_i} = -[Kx_i^2 - p(t) + \Delta]x_i + \xi_0 \sum_{j=1}^N J_{i,j} x_j$$

1. Waidyasooriya, Hasitha, and Masanori Hariyama. "Highly-parallel FPGA accelerator for simulated quantum annealing." *IEEE Transactions on Emerging Topics in Computing* (2019).
2. Goto, Hayato, Kosuke Tatsumura, and Alexander R. Dixon. "Combinatorial optimization by simulating adiabatic bifurcations in nonlinear Hamiltonian systems." *Science advances* 5.4 (2019): eaav2372.

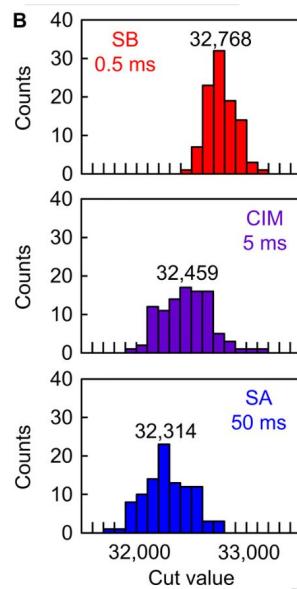
William Larimer Mellon, Founder

Simulated Bifurcation Machine

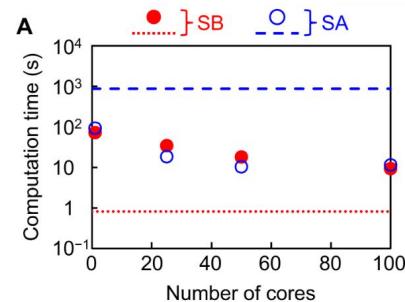
Authors implemented algorithm in FPGA to solve up to 20,000 nodes fully connected graphs



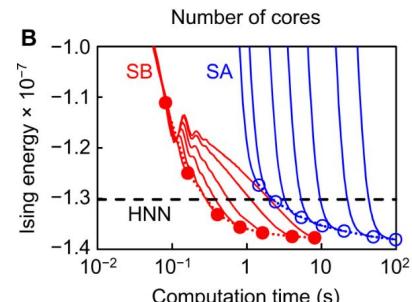
	Ave. of HNN			GW-SDP		
	Best (ms)	Ave. (ms)	Worst (ms)	Best (ms)	Ave. (ms)	Worst (ms)
SB	0.047	0.061	0.074	0.040	0.047	0.058
CIM	0.155	0.769	N/A	0.071	0.264	1.16
SA	2.64	6.80	N/A	2.10	3.20	7.15



Authors implemented algorithm in CPU and GPU to solve up to 1'000,000 nodes fully connected graphs



Dots: CPU
Lines: GPU

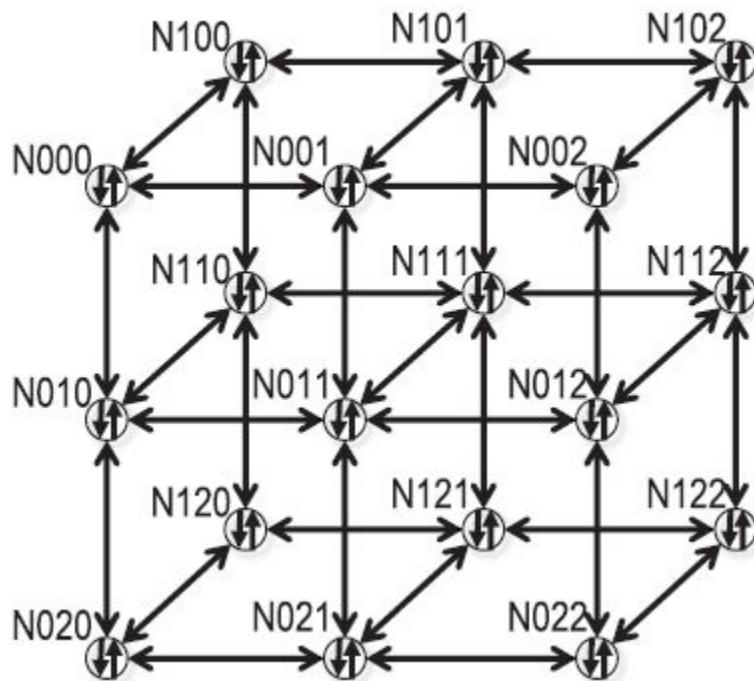


Dots: Cutoff value
Line: 10-run avg.

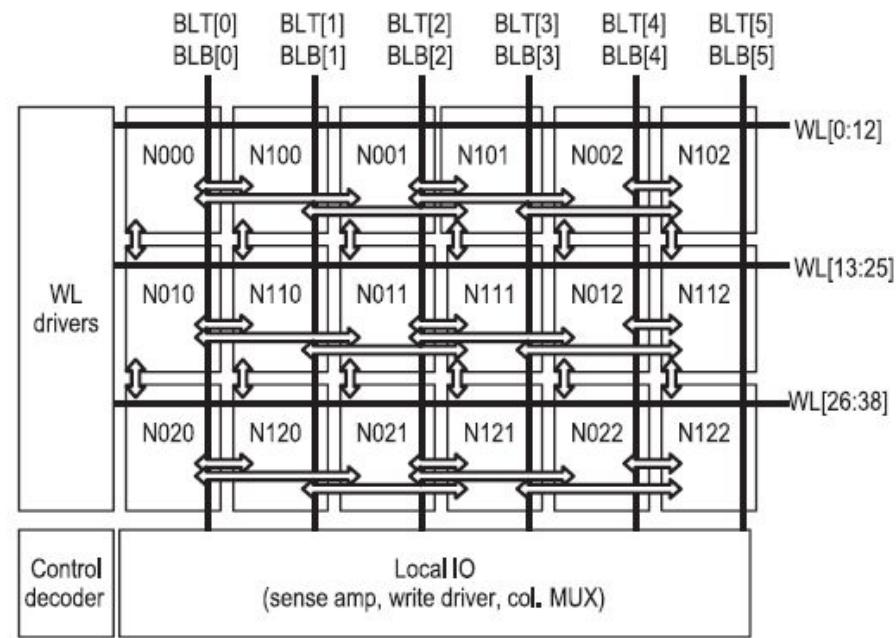
GPU Version Made available through

TOSHIBA aws

Complementary metal-oxide semiconductors (CMOS)

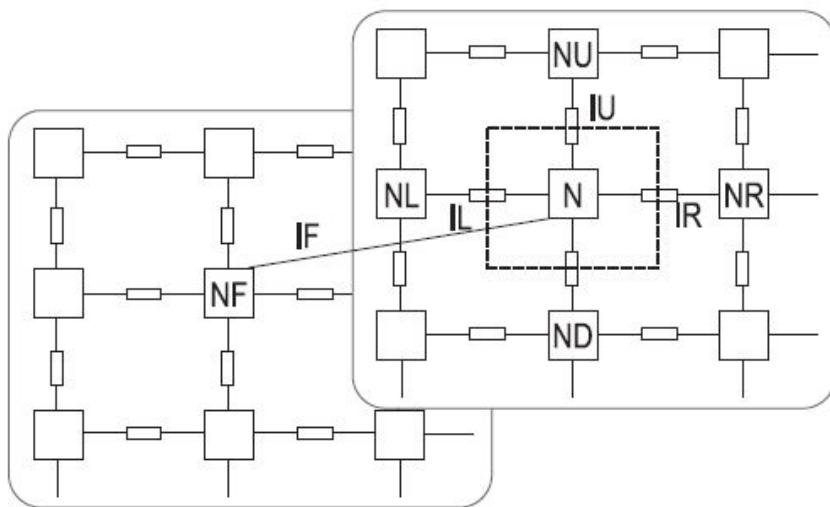


3D Ising Model

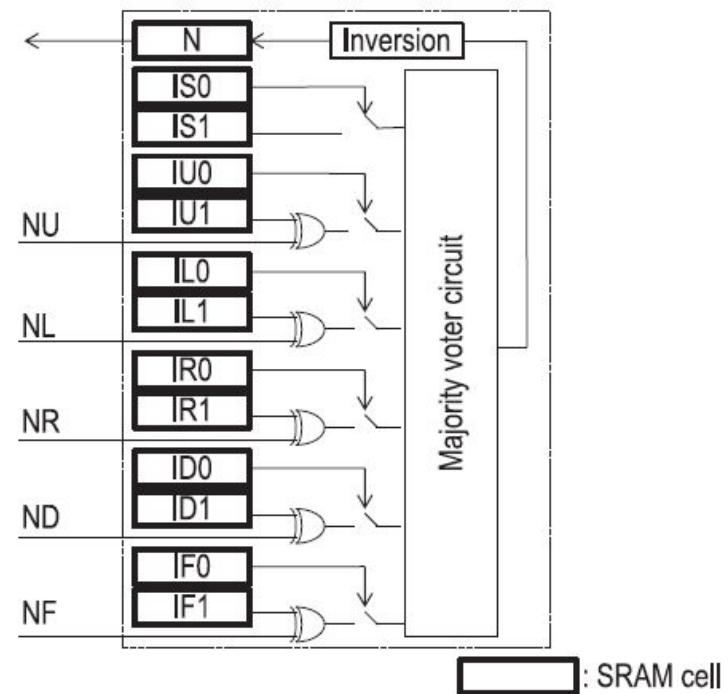


CMOS Static RAM (SRAM) Circuits

Complementary metal-oxide semiconductors (CMOS)



Each Spin has 5 neighbors (Up, Down, Right, Left, Front)



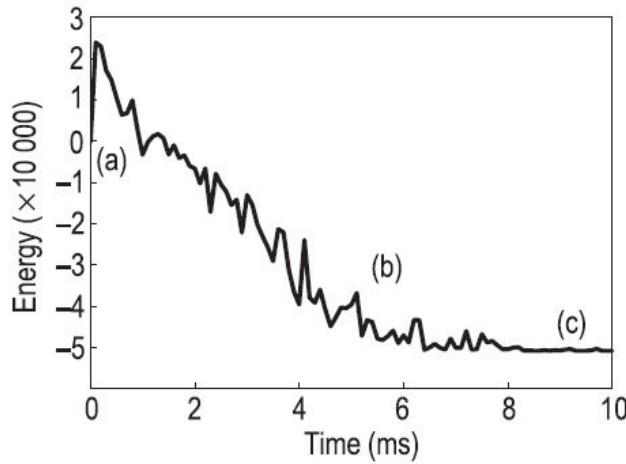
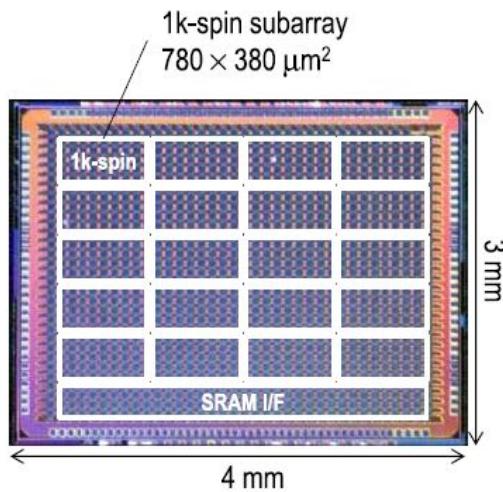
Spin implementation as logic gates

+

Use low voltage to induce random errors in SRAM and jump local minima

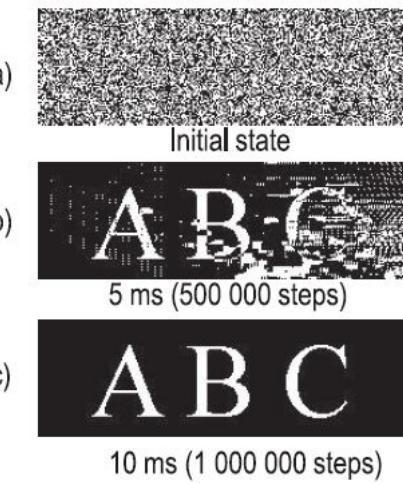
1. Yamaoka, Masanao, et al. "A 20k-spin Ising chip to solve combinatorial optimization problems with CMOS annealing." IEEE Journal of Solid-State Circuits 51.1 (2015): 303-309.

Complementary metal-oxide semiconductors (CMOS)



Actual Chip and Specs

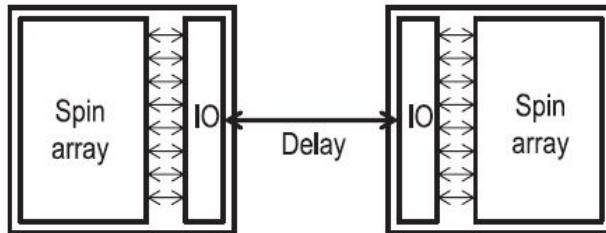
Items	Value
Number of spins	20k (80×256)
Process	65 nm
Chip area	$4 \times 3 = 12 \text{ mm}^2$
Area of spin	$11.27 \times 23.94 = 270 \mu\text{m}^2$
Number of SRAM cells	260k bits Spin value: 1 bit Interaction factor: 2 bit \times 5 = 10 bits External magnetic coefficient: 2 bits
Memory IF	100 MHz
Interaction speed	100 MHz
Operating current of core circuits (1.1 V)	Write: 2.0 mA Read: 6.0 mA Interaction: 44.6 mA



1. Yamaoka, Masanao, et al. "A 20k-spin Ising chip to solve combinatorial optimization problems with CMOS annealing." IEEE Journal of Solid-State Circuits 51.1 (2015): 303-309.

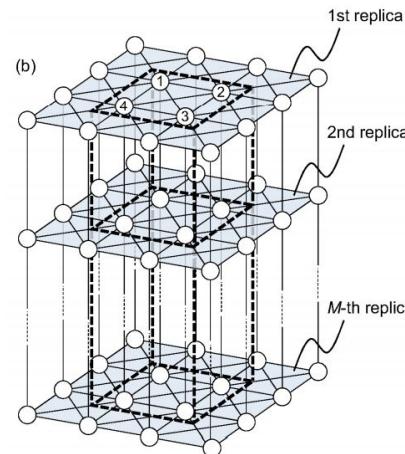
Complementary metal-oxide semiconductors (CMOS)

Easily Parallelizable and manageable

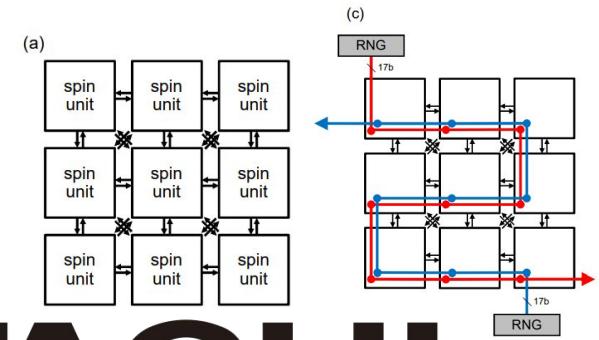


Item	This work
Spin	• SRAM cell (digital bit)
Interaction coefficient & operation	• SRAM cell (digital bit) • Logic circuits (digital bit, asymmetry)
Scalability (one device)	• Easy: CMOS scaling (over 10k spins)
Scalability (multi chip)	• Easy: digital IF can be used
Annealing	• CMOS circuits (digital operation)
Operating condition	• Room temperature (300 K)

Same Idea also implemented in FPGA



King Unit cell useful for both spins and random number generation



Now available through

HITACHI

	Hitachi CMOS Annealer [15], [26]	Hitachi CMOS Annealer [16], [26]
Maximum number of spins Type of coupling Number of couplings Computation Implementation	61,952 King graph 0.37 million not mentioned ASIC	6,400 King graph 0.4 million 8-bit fixed point FPGA

1. Yamaoka, Masanao, et al. "A 20k-spin Ising chip to solve combinatorial optimization problems with CMOS annealing."

2. Okuyama, Takuya, Masato Hayashi, and Masanao Yamaoka. "An Ising computer based on simulated quantum annealing by path integral Monte Carlo method."

Playing with Hitachi's CMOS

Let's go to this interactive interface of the CMOS device from Hitachi

<https://annealing-cloud.com/en/play/ising-editor.html>

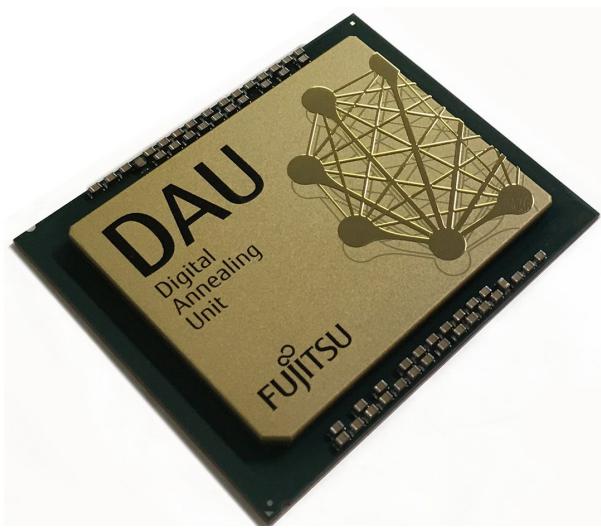
Digital Annealers

CMOS Implementation of Ising solution method

Fully connected 1024 nodes

16-bit precision vs. 4-bit precision D-Wave

*"For obtaining exact solutions of small-size problems,
the SA machine called "Digital Annealer" may be the
fastest so far."*



1. <https://arxiv.org/pdf/1806.08815.pdf>
2. <https://spectrum.ieee.org/tech-talk/computing/hardware/fujitsus-cmos-digital-annealer-produces-quantum-computer-speeds>
3. Goto, Hayato, Kosuke Tatsumura, and Alexander R. Dixon. "Combinatorial optimization by simulating adiabatic bifurcations in nonlinear Hamiltonian systems." *Science advances* 5.4 (2019): eaav2372.

Digital Annealers

Algorithm 1 Simulated Annealing (SA)

```
1: for each run do
2:   initialize to random initial state
3:   for each temperature do
4:     for each MC sweep at this temperature do
5:       for each variable do
6:         propose a flip
7:         if accepted, update the state and effective fields
8:       end for
9:     end for
10:    update the temperature
11:  end for
12: end for
```

Arbitrary start
Initial fields not required
to be calculated

Parallel-trial
Boost acceptance probability

Dynamic offset escape
Helps surmount short, narrow barriers

Algorithm 2 The Digital Annealer's Algorithm

```
1: initial_state  $\leftarrow$  an arbitrary state
2: for each run do
3:   initialize to initial_state
4:    $E_{\text{offset}} \leftarrow 0$ 
5:   for each MC step (iteration) do
6:     if due for temperature update, update the temperature
7:     for each variable  $j$ , in parallel do
8:       propose a flip using  $\Delta E_j - E_{\text{offset}}$ 
9:       if accepted, record
10:      end for
11:      if at least one flip accepted then
12:        choose one flip uniformly at random amongst them
13:        update the state and effective fields, in parallel
14:         $E_{\text{offset}} \leftarrow 0$ 
15:      else
16:         $E_{\text{offset}} \leftarrow E_{\text{offset}} + \text{offset\_increase\_rate}$ 
17:      end if
18:    end for
19:  end for
```

Parallel Tempering

Algorithm 1 Simulated Annealing (SA)

```
1: for each run do
2:   initialize to random initial state
3:   for each temperature do
4:     for each MC sweep at this temperature do
5:       for each variable do
6:         propose a flip
7:         if accepted, update the state and effective fields
8:       end for
9:     end for
10:    update the temperature
11:  end for
12: end for
```

Algorithm 3 Parallel Tempering with Isoenergetic Cluster Moves (PT+ICM)

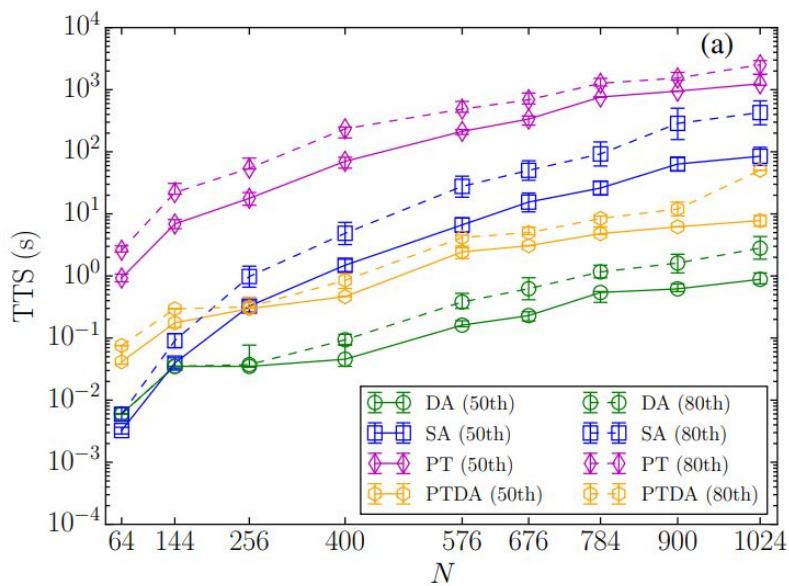
```
1: initialize all replicas with random initial states
2: for each MC sweep do
3:   for each replica, for each variable do
4:     propose a flip
5:     if accepted, update the state and effective fields
6:   end for
7:   for each pair of sequential replicas do
8:     propose a replica exchange
9:     if accepted, swap the temperatures between the replicas
10:   end for
11:   perform ICM update, swapping the states of a cluster of variables that
    have opposite states in the two replicas; update the states and the effective
    fields for both replicas
12: end for
```

- Instead of having a single state you have several replicas
- Then the flips can be done among replicas
- It can be implemented in the Digital Annealer
- Additionally: There can be cluster updates (flip more than one spin if they are “connected”)
 - Similar to Anil’s intuition on the [Swendsen-Wang Algorithm](#)

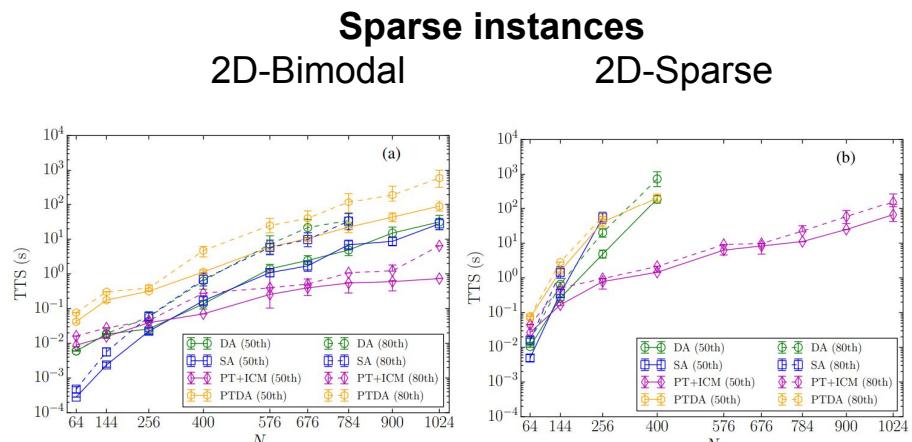
1. <https://arxiv.org/pdf/1806.08815.pdf>
2. <https://spectrum.ieee.org/tech-talk/computing/hardware/fujitsu-cmos-digital-annealer-produces-quantum-computer-speeds>

Digital Annealing v Simulated Annealing v Parallel Tempering

Fully connected instances



Sparse instances



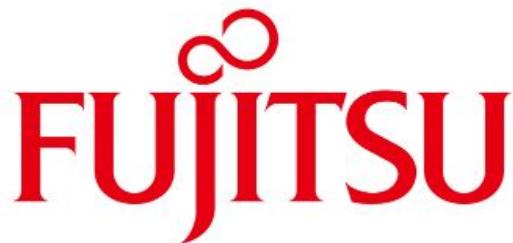
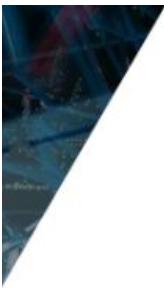
Digital Annealer Wins

- DA Digital Annealer
- SA Simulated Annealing
- PT(+ICM) Parallel Tempering (+Isoenergetic Cluster Moves)
- PTDA Parallel Tempering Digital Annealer

Parallel Tempering Wins



Digital Annealers



Quantum Computing Challenge Series



CH TCO	Quantum Computing Challenge Series - Max Cut Marathon Match Ended Apr 04 Marathon Match	\$11,500 Purse
CH TCO	Quantum Computing Learning Challenge #3 - Max Cut Ended Aug 04 Python Data Science Other	\$250 Purse
CH TCO	Quantum Computing Learning Challenge 2 - Scheduling Ended Feb 28 Python Data Science Other	\$250 Purse
CH TCO	Quantum Computing Learning Challenge #1 - Solve Sudoku Instantly Ended Feb 14 Algorithm Python Data Science +1	\$250 Purse

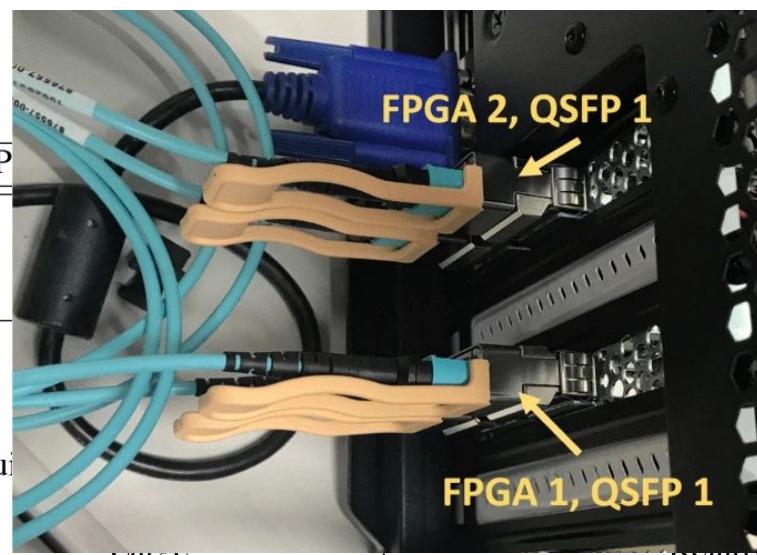


1. <https://arxiv.org/pdf/1806.08815.pdf>
 2. <https://spectrum.ieee.org/tech-talk/computing/hardware/fujitsu-cmos-digital-annealer-produces-quantum-computer-speeds>
 3. <https://tc3-japan.github.io/DA Tutorial/index.html>
- William Larimer Mellon, Founder*

Digital Annealer v Application-specific integrated circuit v FPGA v GPU

	Fujitsu Digital Annealer [25]	Hitachi CMOS Annealer [15], [26]	Hitachi CMOS Annealer [16], [26]	FPGA accelerator
Maximum number of spins	8192	61,952	6,400	32,768
Type of coupling	Total coupling	King graph	King graph	Total Coupling
Number of couplings	67 million	0.37 million	0.4 million	1 billion
Computation	64-bit fixed-point	not mentioned	8-bit fixed point	32-bit floating-point
Implementation	ASIC	ASIC	FPGA	2-FPGA connected via fiber

Category	FP	GPU accelerator
Speed-up	Very large	Implementation
Accuracy	Very small	Iteration
Problem size	From PCs to supercomputers	
Power consumption	Very large	
Power-efficiency	Very small	
Availability	From PCs to supercomputers	
Programmability	OpenCL, OpenAcc, etc.	
Compilation time	Within a minute	
Design time	Relatively small	



- Waidyasooriya, Hasitha Muthumala, and Masanori Hariyama. "A GPU-Based Quantum Annealing Simulator for Fully-Connected Ising Models Utilizing Spatial and Temporal Parallelism."
- Waidyasooriya, H.M., Hariyama, M., Miyama, M.J. et al. OpenCL-based design of an FPGA accelerator for quantum annealing simulation.

³ Waidyasooriya, Hasitha Muthumala, and Masanori Hariyama. "Highly-Parallel FPGA Accelerator for Simulated Quantum Annealing" William Larimer Mellon, Founder

Alternatives available

	Fixstars Optigan	D-Wave 2000Q	Hitachi CMOS Annealing	Fujitsu Digital Annealer	Toshiba SBM
Calculation method	GPU	Quantum annealing	Digital circuit	Digital circuit	GPU
Maximum number of bits	Over 100,000	2,048 (16x16x8)	61,952 (352x176)	1,024 / 8,192	10,000
Coefficient parameter	Digital (32 / 64bit)	Analog (about 5bit)	Digital (3bit)	Digital (16/64 bit)	Digital (32bit)
Combined graph	Fully combined	Chimera graph	King Graph	Fully combined	Fully combined
Total number of combined conversion bits	65,536	64	176	1,024 / 8,192	1,000
API endpoint	Fixstars	D-Wave Cloud	Annealing Cloud Web	DA Cloud	AWS

Playing with Fixstars's Optigan

Let's go to this interactive interface of the GPU implementation from Fixstars Optigan

<https://quantum.fixstars.com/product/demo#demoModal>

Translated version (but you cannot run it)

<https://colab.research.google.com/github/bernalde/QuIP/blob/master/notebooks/Notebook%208%20-%20Amplify%20Tutorials.ipynb>

Closing Remarks

- Non-linear Integer Programs model a variety of real world problems from many domains
- Solving them classically has limitations, especially with non-convex objectives, constrained integer variables
- We explored non-classical approaches based on QUBO/Ising
- GAMA is a general purpose heuristic that utilizes Graver Test-Set
- There are a variety of options to solve Ising model